

01-cifar10-ann

September 6, 2024

```
[1]: import numpy as np
      from keras.models import Sequential
      import matplotlib.pyplot as plt

      from tensorflow import keras
      from keras.layers import Conv2D
      from keras.layers import MaxPooling2D
      from keras.layers import Flatten
      from keras.layers import Dense
      from keras.utils import to_categorical
```

1 Loading CIFAR10 Dataset for ANN Study

```
[2]: cifar=keras.datasets.cifar10
      (X_train,y_train),(X_test,y_test)=cifar.load_data()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170498071/170498071 4s
0us/step

```
[3]: X_train=X_train/255.0
      X_test=X_test/255.0
      n=len(np.unique(y_train))
```

2 2 Hidden Layer Study in ANN

```
[4]: y_train=to_categorical(y_train,n)
      y_test=to_categorical(y_test,n)
      model = Sequential()
      model.add(Flatten(input_shape=(32, 32, 3)))
      model.add(Dense(units=32,activation='relu'))
      model.add(Dense(units=64,activation='relu'))
      model.add(Dense(units=10,activation='softmax'))

      model.compile(optimizer='adam',loss='categorical_crossentropy',
                    metrics=['accuracy'])
```

```
answer=model.fit(X_train,y_train,epochs=10,verbose=2,
                 validation_split=0.2, batch_size=64)
model.evaluate(X_test,y_test)
```

```
/usr/local/lib/python3.10/dist-
packages/keras/src/layers/reshaping/flatten.py:37: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)
```

```
Epoch 1/10
625/625 - 4s - 6ms/step - accuracy: 0.1650 - loss: 2.1417 - val_accuracy: 0.1760
- val_loss: 2.0846
Epoch 2/10
625/625 - 2s - 3ms/step - accuracy: 0.1805 - loss: 2.0732 - val_accuracy: 0.1863
- val_loss: 2.0691
Epoch 3/10
625/625 - 3s - 5ms/step - accuracy: 0.1875 - loss: 2.0638 - val_accuracy: 0.1903
- val_loss: 2.0646
Epoch 4/10
625/625 - 5s - 7ms/step - accuracy: 0.1873 - loss: 2.0599 - val_accuracy: 0.1821
- val_loss: 2.0782
Epoch 5/10
625/625 - 3s - 5ms/step - accuracy: 0.1891 - loss: 2.0557 - val_accuracy: 0.1913
- val_loss: 2.0623
Epoch 6/10
625/625 - 2s - 3ms/step - accuracy: 0.1915 - loss: 2.0561 - val_accuracy: 0.1897
- val_loss: 2.0597
Epoch 7/10
625/625 - 4s - 6ms/step - accuracy: 0.1884 - loss: 2.0531 - val_accuracy: 0.1874
- val_loss: 2.0658
Epoch 8/10
625/625 - 4s - 6ms/step - accuracy: 0.1902 - loss: 2.0479 - val_accuracy: 0.1910
- val_loss: 2.0540
Epoch 9/10
625/625 - 3s - 4ms/step - accuracy: 0.1919 - loss: 2.0499 - val_accuracy: 0.1925
- val_loss: 2.0609
Epoch 10/10
625/625 - 2s - 3ms/step - accuracy: 0.1926 - loss: 2.0505 - val_accuracy: 0.1895
- val_loss: 2.0688
313/313          0s 1ms/step -
accuracy: 0.2020 - loss: 2.0574
```

```
[4]: [2.0551371574401855, 0.19329999387264252]
```

```
[5]: test_loss, test_acc = model.evaluate(X_test, y_test, verbose=2)
print(f'\nTest accuracy for 2 Hidden Layer: {test_acc:.3f}')
```

313/313 - 1s - 2ms/step - accuracy: 0.1933 - loss: 2.0551

Test accuracy for 2 Hidden Layer: 0.193

3 3 Hidden Layer Study in ANN

```
[6]: model1 = Sequential()
model1.add(Flatten(input_shape=(32, 32, 3)))
model1.add(Dense(units=32,activation='relu'))
model1.add(Dense(units=64,activation='relu'))
model1.add(Dense(units=64,activation='relu'))
model1.add(Dense(units=10,activation='softmax'))

model1.compile(optimizer='adam',loss='categorical_crossentropy',
               metrics=['accuracy'])
answer=model1.fit(X_train,y_train,epochs=10,verbose=2,
                 validation_split=0.2,batch_size=64)
```

Epoch 1/10

625/625 - 4s - 6ms/step - accuracy: 0.2867 - loss: 1.9323 - val_accuracy: 0.3414
- val_loss: 1.8155

Epoch 2/10

625/625 - 4s - 7ms/step - accuracy: 0.3614 - loss: 1.7710 - val_accuracy: 0.3784
- val_loss: 1.7356

Epoch 3/10

625/625 - 4s - 6ms/step - accuracy: 0.3812 - loss: 1.7156 - val_accuracy: 0.3831
- val_loss: 1.7230

Epoch 4/10

625/625 - 2s - 4ms/step - accuracy: 0.3987 - loss: 1.6736 - val_accuracy: 0.3878
- val_loss: 1.7219

Epoch 5/10

625/625 - 3s - 4ms/step - accuracy: 0.4070 - loss: 1.6496 - val_accuracy: 0.3884
- val_loss: 1.7010

Epoch 6/10

625/625 - 2s - 3ms/step - accuracy: 0.4140 - loss: 1.6324 - val_accuracy: 0.4141
- val_loss: 1.6514

Epoch 7/10

625/625 - 2s - 3ms/step - accuracy: 0.4199 - loss: 1.6132 - val_accuracy: 0.3879
- val_loss: 1.6883

Epoch 8/10

625/625 - 4s - 6ms/step - accuracy: 0.4206 - loss: 1.6043 - val_accuracy: 0.3906
- val_loss: 1.6876

Epoch 9/10

625/625 - 2s - 3ms/step - accuracy: 0.4261 - loss: 1.5971 - val_accuracy: 0.4138
- val_loss: 1.6224

Epoch 10/10

625/625 - 3s - 4ms/step - accuracy: 0.4300 - loss: 1.5798 - val_accuracy: 0.4070

- val_loss: 1.6751

```
[7]: test_loss, test_acc = model1.evaluate(X_test, y_test, verbose=2)
      print(f'\nTest accuracy for 3 hidden layer: {test_acc:.3f}')
```

313/313 - 0s - 1ms/step - accuracy: 0.4106 - loss: 1.6494

Test accuracy for 3 hidden layer: 0.411

4 1 Hidden Layer Study in ANN

```
[8]: model2 = Sequential()
      model2.add(Flatten(input_shape=(32, 32, 3)))
      model2.add(Dense(units=32,activation='relu'))
      model2.add(Dense(units=10,activation='softmax'))
      model2.compile(optimizer='adam',loss='categorical_crossentropy',
                     metrics=['accuracy'])
      answer=model2.fit(X_train,y_train,epochs=10,verbose=2,
                        validation_split=0.2,batch_size=64)
```

Epoch 1/10

625/625 - 3s - 5ms/step - accuracy: 0.2573 - loss: 2.0219 - val_accuracy: 0.2939
- val_loss: 1.9700

Epoch 2/10

625/625 - 4s - 6ms/step - accuracy: 0.3056 - loss: 1.9076 - val_accuracy: 0.3091
- val_loss: 1.9057

Epoch 3/10

625/625 - 4s - 6ms/step - accuracy: 0.3198 - loss: 1.8722 - val_accuracy: 0.3197
- val_loss: 1.8838

Epoch 4/10

625/625 - 3s - 4ms/step - accuracy: 0.3205 - loss: 1.8584 - val_accuracy: 0.3151
- val_loss: 1.8869

Epoch 5/10

625/625 - 2s - 3ms/step - accuracy: 0.3301 - loss: 1.8443 - val_accuracy: 0.3292
- val_loss: 1.8511

Epoch 6/10

625/625 - 4s - 6ms/step - accuracy: 0.3334 - loss: 1.8356 - val_accuracy: 0.3187
- val_loss: 1.8665

Epoch 7/10

625/625 - 5s - 7ms/step - accuracy: 0.3346 - loss: 1.8285 - val_accuracy: 0.3320
- val_loss: 1.8433

Epoch 8/10

625/625 - 5s - 7ms/step - accuracy: 0.3390 - loss: 1.8243 - val_accuracy: 0.3417
- val_loss: 1.8367

Epoch 9/10

625/625 - 3s - 5ms/step - accuracy: 0.3401 - loss: 1.8200 - val_accuracy: 0.3353
- val_loss: 1.8498

Epoch 10/10

625/625 - 3s - 4ms/step - accuracy: 0.3415 - loss: 1.8133 - val_accuracy: 0.3386
- val_loss: 1.8288

```
[9]: test_loss, test_acc = model2.evaluate(X_test, y_test, verbose=2)
      print(f'\nTest accuracy for 1 hidden layer: {test_acc:.3f}')
```

313/313 - 0s - 1ms/step - accuracy: 0.3374 - loss: 1.8132

Test accuracy for 1 hidden layer: 0.337

5 Conclusion

As we decrease the number of hidden layers, the test accuracy decreases due to limited learning capacity, underfitting and loss of hierarchical feature extraction.

CIFAR10 Dataset gives less accuracy than FashionMNIST Dataset, which means that CIFAR10 is more complex than FashionMNIST. CIFAR10 requires more number of CNN Blocks to reach the accuracy of that of the FashionMNIST.

02-cifar10-cnn

September 6, 2024

```
[1]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score, classification_report

import keras
import tensorflow as tf
from tensorflow import keras
from keras.models import Sequential
from tensorflow.keras import regularizers, optimizers
from tensorflow.keras.layers import Input, Conv2D, Dense, Flatten, Dropout
from tensorflow.keras.layers import GlobalMaxPooling2D, MaxPooling2D
from tensorflow.keras.layers import BatchNormalization
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.models import Model
```

1 Loading CIFAR10 DataSet for CNN Study

```
[2]: from tensorflow.keras.datasets import cifar10
(X_train, Y_train), (X_test, Y_test) = cifar10.load_data()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170498071/170498071 2s
0us/step

```
[3]: X_train = X_train/255
X_test = X_test/255
Y_train_en = to_categorical(Y_train,10)
Y_test_en = to_categorical(Y_test,10)
```

2 3 Block CNN Study

```
[4]: model = Sequential()
model.add(Conv2D(32,(4,4),input_shape = (32,32,3),activation='relu'))
model.add(MaxPooling2D(pool_size = (2,2)))
model.add(Conv2D(32,(4,4),input_shape = (32,32,3),activation='relu'))
model.add(MaxPooling2D(pool_size = (2,2)))
```

```

model.add(Conv2D(32,(4,4),input_shape = (32,32,3),activation='relu'))
model.add(MaxPooling2D(pool_size = (2,2)))

model.add(Flatten())
model.add(Dense(128, activation = 'relu'))
model.add(Dense(10, activation = 'softmax'))
model.compile(loss = 'categorical_crossentropy', optimizer = 'adam',
metrics = ['accuracy'])

model.summary()
history = model.fit(X_train, Y_train_en, epochs = 10, verbose=1,
                    validation_data=(X_test,Y_test_en))

```

```

/usr/local/lib/python3.10/dist-
packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not
pass an `input_shape`/`input_dim` argument to a layer. When using Sequential
models, prefer using an `Input(shape)` object as the first layer in the model
instead.

```

```

super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```

Model: "sequential"

| Layer (type) | Output Shape | |
|--------------------------------|--------------------|--|
| Param # | | |
| conv2d (Conv2D) | (None, 29, 29, 32) | |
| ↳1,568 | | |
| max_pooling2d (MaxPooling2D) | (None, 14, 14, 32) | |
| ↳ 0 | | |
| conv2d_1 (Conv2D) | (None, 11, 11, 32) | |
| ↳16,416 | | |
| max_pooling2d_1 (MaxPooling2D) | (None, 5, 5, 32) | |
| ↳ 0 | | |
| conv2d_2 (Conv2D) | (None, 2, 2, 32) | |
| ↳16,416 | | |
| max_pooling2d_2 (MaxPooling2D) | (None, 1, 1, 32) | |
| ↳ 0 | | |
| flatten (Flatten) | (None, 32) | |
| ↳ 0 | | |

dense (Dense) (None, 128) └
└4,224

dense_1 (Dense) (None, 10) └
└1,290

Total params: 39,914 (155.91 KB)

Trainable params: 39,914 (155.91 KB)

Non-trainable params: 0 (0.00 B)

Epoch 1/10

1563/1563 68s 42ms/step -
accuracy: 0.3132 - loss: 1.8318 - val_accuracy: 0.5151 - val_loss: 1.3419

Epoch 2/10

1563/1563 63s 40ms/step -
accuracy: 0.5172 - loss: 1.3372 - val_accuracy: 0.5497 - val_loss: 1.2618

Epoch 3/10

1563/1563 82s 40ms/step -
accuracy: 0.5743 - loss: 1.1963 - val_accuracy: 0.5842 - val_loss: 1.1662

Epoch 4/10

1563/1563 84s 41ms/step -
accuracy: 0.6081 - loss: 1.1027 - val_accuracy: 0.5873 - val_loss: 1.1687

Epoch 5/10

1563/1563 81s 41ms/step -
accuracy: 0.6371 - loss: 1.0317 - val_accuracy: 0.6256 - val_loss: 1.0631

Epoch 6/10

1563/1563 84s 42ms/step -
accuracy: 0.6553 - loss: 0.9804 - val_accuracy: 0.6397 - val_loss: 1.0304

Epoch 7/10

1563/1563 79s 40ms/step -
accuracy: 0.6691 - loss: 0.9407 - val_accuracy: 0.6416 - val_loss: 1.0386

Epoch 8/10

1563/1563 83s 41ms/step -
accuracy: 0.6824 - loss: 0.9031 - val_accuracy: 0.6287 - val_loss: 1.0610

Epoch 9/10

1563/1563 83s 42ms/step -
accuracy: 0.6987 - loss: 0.8592 - val_accuracy: 0.6595 - val_loss: 0.9947

Epoch 10/10

1563/1563 82s 42ms/step -
accuracy: 0.7021 - loss: 0.8377 - val_accuracy: 0.6596 - val_loss: 0.9881


```
[5]: test_loss, test_acc = model.evaluate(X_test, Y_test_en, verbose=2)
      print(f'\nTest Accuracy 3 Block Study: {test_acc:.3f}')
```

313/313 - 3s - 10ms/step - accuracy: 0.6596 - loss: 0.9881

Test Accuracy 3 Block Study: 0.660

3 2 Blocks CNN Study

```
[6]: model1 = Sequential()
      model1.add(Conv2D(32,(4,4),input_shape = (32,32,3),activation='relu'))
      model1.add(MaxPooling2D(pool_size = (2,2)))
      model1.add(Conv2D(32,(4,4),input_shape = (32,32,3),activation='relu'))
      model1.add(MaxPooling2D(pool_size = (2,2)))

      model1.add(Flatten())
      model1.add(Dense(128, activation = 'relu'))
      model1.add(Dense(10, activation = 'softmax'))
      model1.compile(loss = 'categorical_crossentropy', optimizer = 'adam',
                     metrics = ['accuracy'])

      model1.summary()
      history = model1.fit(X_train, Y_train_en, epochs = 10, verbose=1,
                           validation_data=(X_test,Y_test_en))
```

Model: "sequential_1"

| Layer (type) ↳Param # | Output Shape | |
|---------------------------------------|--------------------|---|
| conv2d_3 (Conv2D) ↳1,568 | (None, 29, 29, 32) | ↳ |
| max_pooling2d_3 (MaxPooling2D) ↳ 0 | (None, 14, 14, 32) | ↳ |
| conv2d_4 (Conv2D) ↳16,416 | (None, 11, 11, 32) | ↳ |
| max_pooling2d_4 (MaxPooling2D) ↳ 0 | (None, 5, 5, 32) | ↳ |
| flatten_1 (Flatten) ↳ 0 | (None, 800) | ↳ |

```
dense_2 (Dense)                (None, 128)
↳102,528

dense_3 (Dense)                (None, 10)
↳1,290
```

Total params: 121,802 (475.79 KB)

Trainable params: 121,802 (475.79 KB)

Non-trainable params: 0 (0.00 B)

```
Epoch 1/10
1563/1563          71s 44ms/step -
accuracy: 0.3812 - loss: 1.6963 - val_accuracy: 0.5346 - val_loss: 1.2978
Epoch 2/10
1563/1563          78s 42ms/step -
accuracy: 0.5625 - loss: 1.2307 - val_accuracy: 0.5919 - val_loss: 1.1364
Epoch 3/10
1563/1563          83s 43ms/step -
accuracy: 0.6260 - loss: 1.0706 - val_accuracy: 0.6197 - val_loss: 1.0800
Epoch 4/10
1563/1563          83s 43ms/step -
accuracy: 0.6563 - loss: 0.9803 - val_accuracy: 0.6440 - val_loss: 1.0192
Epoch 5/10
1563/1563          79s 41ms/step -
accuracy: 0.6848 - loss: 0.9058 - val_accuracy: 0.6553 - val_loss: 0.9969
Epoch 6/10
1563/1563          81s 40ms/step -
accuracy: 0.7089 - loss: 0.8364 - val_accuracy: 0.6530 - val_loss: 1.0215
Epoch 7/10
1563/1563          67s 43ms/step -
accuracy: 0.7305 - loss: 0.7706 - val_accuracy: 0.6726 - val_loss: 0.9719
Epoch 8/10
1563/1563          82s 43ms/step -
accuracy: 0.7498 - loss: 0.7188 - val_accuracy: 0.6767 - val_loss: 0.9672
Epoch 9/10
1563/1563          63s 40ms/step -
accuracy: 0.7664 - loss: 0.6682 - val_accuracy: 0.6817 - val_loss: 0.9717
Epoch 10/10
1563/1563          67s 43ms/step -
accuracy: 0.7791 - loss: 0.6231 - val_accuracy: 0.6703 - val_loss: 1.0324
```

```
[7]: test_loss, test_acc = model1.evaluate(X_test, Y_test_en, verbose=2)
print(f'\nTest accuracy 2 Block Study: {test_acc:.3f}')
```

313/313 - 3s - 10ms/step - accuracy: 0.6596 - loss: 0.9881

Test accuracy 2 Block Study: 0.660

4 1 Block CNN Study

```
[9]: model2 = Sequential()
model2.add(Conv2D(32,(4,4),input_shape = (32,32,3),activation='relu'))
model2.add(MaxPooling2D(pool_size = (2,2)))

model2.add(Flatten())
model2.add(Dense(128, activation = 'relu'))
model2.add(Dense(10, activation = 'softmax'))
model2.compile(loss = 'categorical_crossentropy', optimizer = 'adam',
metrics = ['accuracy'])

model2.summary()
history = model2.fit(X_train, Y_train_en, epochs = 10,
verbose=1,validation_data=(X_test,Y_test_en))
```

/usr/local/lib/python3.10/dist-

packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Model: "sequential_3"

| Layer (type) | Output Shape | |
|--------------------------------|--------------------|--|
| Param # | | |
| conv2d_6 (Conv2D) | (None, 29, 29, 32) | |
| ↳1,568 | | |
| max_pooling2d_6 (MaxPooling2D) | (None, 14, 14, 32) | |
| ↳ 0 | | |
| flatten_3 (Flatten) | (None, 6272) | |
| ↳ 0 | | |
| dense_6 (Dense) | (None, 128) | |
| ↳802,944 | | |

dense_7 (Dense) (None, 10)
↩1,290

Total params: 805,802 (3.07 MB)

Trainable params: 805,802 (3.07 MB)

Non-trainable params: 0 (0.00 B)

Epoch 1/10
1563/1563 57s 36ms/step -
accuracy: 0.4033 - loss: 1.6532 - val_accuracy: 0.5592 - val_loss: 1.2493
Epoch 2/10
1563/1563 81s 35ms/step -
accuracy: 0.5867 - loss: 1.1764 - val_accuracy: 0.6082 - val_loss: 1.1088
Epoch 3/10
1563/1563 82s 35ms/step -
accuracy: 0.6439 - loss: 1.0229 - val_accuracy: 0.6273 - val_loss: 1.0552
Epoch 4/10
1563/1563 56s 36ms/step -
accuracy: 0.6751 - loss: 0.9319 - val_accuracy: 0.6416 - val_loss: 1.0583
Epoch 5/10
1563/1563 82s 35ms/step -
accuracy: 0.7084 - loss: 0.8370 - val_accuracy: 0.6358 - val_loss: 1.0507
Epoch 6/10
1563/1563 82s 36ms/step -
accuracy: 0.7390 - loss: 0.7568 - val_accuracy: 0.6433 - val_loss: 1.0747
Epoch 7/10
1563/1563 82s 36ms/step -
accuracy: 0.7560 - loss: 0.7044 - val_accuracy: 0.6329 - val_loss: 1.0941
Epoch 8/10
1563/1563 82s 35ms/step -
accuracy: 0.7787 - loss: 0.6370 - val_accuracy: 0.6370 - val_loss: 1.1502
Epoch 9/10
1563/1563 60s 38ms/step -
accuracy: 0.7974 - loss: 0.5741 - val_accuracy: 0.6451 - val_loss: 1.1211
Epoch 10/10
1563/1563 78s 36ms/step -
accuracy: 0.8219 - loss: 0.5147 - val_accuracy: 0.6463 - val_loss: 1.1989

```
[10]: test_loss, test_acc = model2.evaluate(X_test, Y_test_en, verbose=2)
      print(f'\nTest accuracy for 1 Block Study: {test_acc:.3f}')
```

313/313 - 2s - 8ms/step - accuracy: 0.6463 - loss: 1.1989

Test accuracy for 1 Block Study: 0.646

5 Conclusion (CIFAR 10 CNN Study):

If your problem requires detailed feature extraction, decreasing the number of blocks too much may decrease the performance.

In 2 Block and 3 Block CNN Study, accuracy (0.66) didn't changed but in 1 Block CNN, it got decreased to (0.646).

CIFAR10 shows lesser accuracy than FashionMNIST at the same no. of hidden layers, which means CIFAR10 is more complex dataset than FashionMNIST. CIFAR 10 requires more no. of hidden layers to reach the accuracy level of that of the FashionMNIST.

03-fashionmnist-ann

September 6, 2024

```
[1]: import matplotlib.pyplot as plt
import numpy as np

from tensorflow import keras
from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
from keras.models import Sequential
from keras.utils import to_categorical
```

1 Loading FashionMNIST Dataset

```
[2]: fashion_mnist = keras.datasets.fashion_mnist
(X_train,y_train),(X_test,y_test) = fashion_mnist.load_data()
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/train-labels-idx1-ubyte.gz
29515/29515          0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/train-images-idx3-ubyte.gz
26421880/26421880    2s
0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/t10k-labels-idx1-ubyte.gz
5148/5148            0s 1us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/t10k-images-idx3-ubyte.gz
4422102/4422102      1s
0us/step
```

```
[3]: X_train=X_train/255.0
X_test=X_test/255.0
n=len(np.unique(y_train))
```

```
[4]: y_train=to_categorical(y_train,n)
y_test=to_categorical(y_test,n)
```

2 2 Hidden Layer ANN

```
[5]: model = Sequential()
model.add(Flatten(input_shape=(28, 28)))
model.add(Dense(units=32,activation='relu'))
model.add(Dense(units=64,activation='relu'))
model.add(Dense(units=10,activation='softmax'))
model.compile(optimizer='adam',loss='categorical_crossentropy',
              metrics=['accuracy'])
answer=model.fit(X_train,y_train,epochs=10,verbose=2,
                validation_split=0.2,batch_size=64)
```

```
/usr/local/lib/python3.10/dist-
packages/keras/src/layers/reshaping/flatten.py:37: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)
```

Epoch 1/10

750/750 - 4s - 6ms/step - accuracy: 0.7916 - loss: 0.5993 - val_accuracy: 0.8402
- val_loss: 0.4618

Epoch 2/10

750/750 - 2s - 2ms/step - accuracy: 0.8513 - loss: 0.4201 - val_accuracy: 0.8531
- val_loss: 0.4127

Epoch 3/10

750/750 - 3s - 3ms/step - accuracy: 0.8632 - loss: 0.3787 - val_accuracy: 0.8610
- val_loss: 0.3865

Epoch 4/10

750/750 - 2s - 3ms/step - accuracy: 0.8719 - loss: 0.3561 - val_accuracy: 0.8702
- val_loss: 0.3660

Epoch 5/10

750/750 - 3s - 5ms/step - accuracy: 0.8784 - loss: 0.3360 - val_accuracy: 0.8734
- val_loss: 0.3532

Epoch 6/10

750/750 - 2s - 3ms/step - accuracy: 0.8818 - loss: 0.3216 - val_accuracy: 0.8703
- val_loss: 0.3602

Epoch 7/10

750/750 - 2s - 2ms/step - accuracy: 0.8853 - loss: 0.3117 - val_accuracy: 0.8725
- val_loss: 0.3596

Epoch 8/10

750/750 - 2s - 3ms/step - accuracy: 0.8866 - loss: 0.3055 - val_accuracy: 0.8716
- val_loss: 0.3533

Epoch 9/10

750/750 - 2s - 3ms/step - accuracy: 0.8925 - loss: 0.2918 - val_accuracy: 0.8780
- val_loss: 0.3371

Epoch 10/10

750/750 - 2s - 2ms/step - accuracy: 0.8942 - loss: 0.2860 - val_accuracy: 0.8767
- val_loss: 0.3536

```
[6]: test_loss, test_acc = model.evaluate(X_test, y_test, verbose=2)
      print(f'\nTest accuracy for 2 Hidden Layers: {test_acc:.3f}')
```

313/313 - 1s - 2ms/step - accuracy: 0.8677 - loss: 0.3839

Test accuracy for 2 Hidden Layers: 0.868

3 3 hidden layer ANN

```
[7]: model1 = Sequential()
      model1.add(Flatten(input_shape=(28, 28)))
      model1.add(Dense(units=32,activation='relu'))
      model1.add(Dense(units=64,activation='relu'))
      model1.add(Dense(units=64,activation='relu'))
      model1.add(Dense(units=10,activation='softmax'))
      model1.compile(optimizer='adam',loss='categorical_crossentropy',
                     metrics=['accuracy'])
      answer=model1.fit(X_train,y_train,epochs=10,verbose=2,
                        validation_split=0.2,batch_size=64)
```

Epoch 1/10

750/750 - 4s - 5ms/step - accuracy: 0.7864 - loss: 0.6090 - val_accuracy: 0.8408
- val_loss: 0.4511

Epoch 2/10

750/750 - 2s - 3ms/step - accuracy: 0.8515 - loss: 0.4148 - val_accuracy: 0.8538
- val_loss: 0.4029

Epoch 3/10

750/750 - 3s - 3ms/step - accuracy: 0.8643 - loss: 0.3739 - val_accuracy: 0.8478
- val_loss: 0.4099

Epoch 4/10

750/750 - 2s - 3ms/step - accuracy: 0.8718 - loss: 0.3508 - val_accuracy: 0.8656
- val_loss: 0.3736

Epoch 5/10

750/750 - 3s - 4ms/step - accuracy: 0.8786 - loss: 0.3328 - val_accuracy: 0.8707
- val_loss: 0.3574

Epoch 6/10

750/750 - 2s - 2ms/step - accuracy: 0.8813 - loss: 0.3189 - val_accuracy: 0.8728
- val_loss: 0.3513

Epoch 7/10

750/750 - 3s - 4ms/step - accuracy: 0.8864 - loss: 0.3097 - val_accuracy: 0.8758
- val_loss: 0.3435

Epoch 8/10

750/750 - 2s - 2ms/step - accuracy: 0.8900 - loss: 0.2977 - val_accuracy: 0.8732
- val_loss: 0.3496

Epoch 9/10

750/750 - 2s - 2ms/step - accuracy: 0.8932 - loss: 0.2872 - val_accuracy: 0.8798
- val_loss: 0.3336

Epoch 10/10
750/750 - 3s - 4ms/step - accuracy: 0.8958 - loss: 0.2807 - val_accuracy: 0.8731
- val_loss: 0.3494

```
[8]: test_loss, test_acc = model1.evaluate(X_test, y_test, verbose=2)
      print(f'\nTest accuracy for 3 Hidden Layers: {test_acc:.3f}')
```

313/313 - 1s - 2ms/step - accuracy: 0.8653 - loss: 0.3724

Test accuracy for 3 Hidden Layers: 0.865

4 1 Hidden Layer ANN

```
[9]: model2 = Sequential()
      model2.add(Flatten(input_shape=(28,28)))
      model2.add(Dense(units=32,activation='relu'))
      model2.add(Dense(units=10,activation='softmax'))
      model2.compile(optimizer='adam',loss='categorical_crossentropy',
                     metrics=['accuracy'])
      answer=model2.fit(X_train,y_train,epochs=10,verbose=2,
                       validation_split=0.2,batch_size=64)
```

Epoch 1/10
750/750 - 3s - 3ms/step - accuracy: 0.7848 - loss: 0.6347 - val_accuracy: 0.8336
- val_loss: 0.4684

Epoch 2/10
750/750 - 2s - 3ms/step - accuracy: 0.8467 - loss: 0.4413 - val_accuracy: 0.8467
- val_loss: 0.4430

Epoch 3/10
750/750 - 1s - 2ms/step - accuracy: 0.8571 - loss: 0.4067 - val_accuracy: 0.8583
- val_loss: 0.4096

Epoch 4/10
750/750 - 2s - 3ms/step - accuracy: 0.8646 - loss: 0.3852 - val_accuracy: 0.8577
- val_loss: 0.4093

Epoch 5/10
750/750 - 4s - 5ms/step - accuracy: 0.8689 - loss: 0.3691 - val_accuracy: 0.8608
- val_loss: 0.3904

Epoch 6/10
750/750 - 2s - 2ms/step - accuracy: 0.8715 - loss: 0.3587 - val_accuracy: 0.8652
- val_loss: 0.3805

Epoch 7/10
750/750 - 3s - 3ms/step - accuracy: 0.8769 - loss: 0.3469 - val_accuracy: 0.8703
- val_loss: 0.3689

Epoch 8/10
750/750 - 2s - 3ms/step - accuracy: 0.8786 - loss: 0.3394 - val_accuracy: 0.8634
- val_loss: 0.3815

Epoch 9/10

```
750/750 - 3s - 4ms/step - accuracy: 0.8808 - loss: 0.3317 - val_accuracy: 0.8627  
- val_loss: 0.3890  
Epoch 10/10  
750/750 - 4s - 5ms/step - accuracy: 0.8828 - loss: 0.3254 - val_accuracy: 0.8652  
- val_loss: 0.3700
```

```
[11]: test_loss, test_acc = model2.evaluate(X_test, y_test, verbose=2)  
      print(f'\nTest accuracy for 1 Hidden Layer: {test_acc:.3f}')
```

```
313/313 - 0s - 1ms/step - accuracy: 0.8610 - loss: 0.3981
```

```
Test accuracy for 1 Hidden Layer: 0.861
```

5 Conclusion

As the FashionMNIST Dataset is simpler than CIFAR10 Dataset, it shows more test accuracy at the same no. of hidden layer than that in CIFAR10.

For 2 hidden layers it was the highest means than that of 1 and 3 hidden layer architectures. Test Accuracy was lower for 1 hidden layer architecture than that of 3 layer architecture.

As we decrease the hidden layer, test accuracy decreases (at 1 hidden layer architecture) due to limited learning capacity, underfitting and loss of hierarchial feature extraction. It also suffers overfitting (at 3 hidden layer architecture) which lead to decrement in test accuracy.

04-fashionmnist-cnn

September 6, 2024

```
[1]: from sklearn.metrics import classification_report
import numpy as np
import matplotlib.pyplot as plt

import keras
import tensorflow as tf
from tensorflow import keras
from keras.models import Sequential
from tensorflow.keras import regularizers, optimizers
from tensorflow.keras.layers import Input, Conv2D, Dense, Flatten, Dropout
from tensorflow.keras.layers import GlobalMaxPooling2D, MaxPooling2D
from tensorflow.keras.layers import BatchNormalization
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.models import Model
```

1 Loading FashionMNIST Dataset

```
[2]: (x_train, y_train), (x_test, y_test) = keras.datasets.fashion_mnist.load_data()
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/train-labels-idx1-ubyte.gz
29515/29515          0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/train-images-idx3-ubyte.gz
26421880/26421880    0s
0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/t10k-labels-idx1-ubyte.gz
5148/5148           0s 1us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/t10k-images-idx3-ubyte.gz
4422102/4422102      0s
0us/step
```

```
[3]: x_train = x_train / 255.0
x_test = x_test / 255.0
```

2 3 Block CNN

```
[4]: model = Sequential()
model.add(Conv2D(filters=32, kernel_size=(3, 3), strides=(1, 1),
                 padding='valid', activation='relu',
                 input_shape=(28, 28, 1)))
model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
model.add(Dropout(rate=0.25))

model.add(Conv2D(filters=32, kernel_size=(3, 3),strides=(1, 1),
                 padding='valid', activation='relu',
                 input_shape=(28, 28, 1)))
model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
model.add(Dropout(rate=0.25))

model.add(Conv2D(filters=32, kernel_size=(3, 3),strides=(1, 1),
                 padding='valid', activation='relu', input_shape=(28, 28, 1)))
model.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
model.add(Dropout(rate=0.25))

model.add(Flatten())
model.add(Dense(units=128, activation='relu'))
model.add(Dense(units=10, activation='softmax'))
model.compile(loss=keras.losses.sparse_categorical_crossentropy,
              optimizer=keras.optimizers.Adam(), metrics=['accuracy'])
model.summary()
```

```
/usr/local/lib/python3.10/dist-
packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not
pass an `input_shape`/`input_dim` argument to a layer. When using Sequential
models, prefer using an `Input(shape)` object as the first layer in the model
instead.
```

```
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
Model: "sequential"
```

| Layer (type) | Output Shape | |
|------------------------------|--------------------|---|
| ↳ Param # | | |
| conv2d (Conv2D) | (None, 26, 26, 32) | ↳ |
| ↳ 320 | | |
| max_pooling2d (MaxPooling2D) | (None, 13, 13, 32) | ↳ |
| ↳ 0 | | |

| | | |
|--------------------------------|--------------------|---|
| dropout (Dropout) | (None, 13, 13, 32) | └ |
| ↪ 0 | | |
| conv2d_1 (Conv2D) | (None, 11, 11, 32) | └ |
| ↪ 9,248 | | |
| max_pooling2d_1 (MaxPooling2D) | (None, 5, 5, 32) | └ |
| ↪ 0 | | |
| dropout_1 (Dropout) | (None, 5, 5, 32) | └ |
| ↪ 0 | | |
| conv2d_2 (Conv2D) | (None, 3, 3, 32) | └ |
| ↪ 9,248 | | |
| max_pooling2d_2 (MaxPooling2D) | (None, 1, 1, 32) | └ |
| ↪ 0 | | |
| dropout_2 (Dropout) | (None, 1, 1, 32) | └ |
| ↪ 0 | | |
| flatten (Flatten) | (None, 32) | └ |
| ↪ 0 | | |
| dense (Dense) | (None, 128) | └ |
| ↪ 4,224 | | |
| dense_1 (Dense) | (None, 10) | └ |
| ↪ 1,290 | | |

Total params: 24,330 (95.04 KB)

Trainable params: 24,330 (95.04 KB)

Non-trainable params: 0 (0.00 B)

```
[5]: x_train = np.expand_dims(x_train, -1)
      x_test = np.expand_dims(x_test, -1)
      history = model.fit(x_train, y_train, batch_size=256, epochs=10,
                          validation_split=0.2, verbose=1)
```

Epoch 1/10

188/188

42s 209ms/step -

```

accuracy: 0.3675 - loss: 1.6517 - val_accuracy: 0.7321 - val_loss: 0.7728
Epoch 2/10
188/188          39s 209ms/step -
accuracy: 0.6906 - loss: 0.8311 - val_accuracy: 0.7638 - val_loss: 0.6302
Epoch 3/10
188/188          41s 210ms/step -
accuracy: 0.7338 - loss: 0.7217 - val_accuracy: 0.7922 - val_loss: 0.5814
Epoch 4/10
188/188          36s 193ms/step -
accuracy: 0.7548 - loss: 0.6669 - val_accuracy: 0.8026 - val_loss: 0.5480
Epoch 5/10
188/188          43s 203ms/step -
accuracy: 0.7693 - loss: 0.6294 - val_accuracy: 0.8145 - val_loss: 0.5170
Epoch 6/10
188/188          40s 199ms/step -
accuracy: 0.7833 - loss: 0.5985 - val_accuracy: 0.8163 - val_loss: 0.5087
Epoch 7/10
188/188          39s 191ms/step -
accuracy: 0.7852 - loss: 0.5799 - val_accuracy: 0.8285 - val_loss: 0.4807
Epoch 8/10
188/188          37s 199ms/step -
accuracy: 0.7963 - loss: 0.5527 - val_accuracy: 0.8300 - val_loss: 0.4736
Epoch 9/10
188/188          40s 211ms/step -
accuracy: 0.8082 - loss: 0.5326 - val_accuracy: 0.8344 - val_loss: 0.4570
Epoch 10/10
188/188          36s 193ms/step -
accuracy: 0.8103 - loss: 0.5216 - val_accuracy: 0.8446 - val_loss: 0.4335

```

```

[6]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
      print(f'\nTest accuracy for 3 block CNN: {test_acc:.3f}')

```

```

313/313 - 2s - 8ms/step - accuracy: 0.8442 - loss: 0.4446

```

```

Test accuracy for 3 block CNN: 0.844

```

3 2 Blocks CNN

```

[7]: model1 = Sequential()
      model1.add(Conv2D(filters=32, kernel_size=(3, 3), strides=(1, 1),
                        padding='valid', activation='relu', input_shape=(28, 28, 1)))
      model1.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
      model1.add(Dropout(rate=0.25))

      model1.add(Conv2D(filters=32, kernel_size=(3, 3), strides=(1, 1),
                        padding='valid', activation='relu', input_shape=(28, 28, 1)))
      model1.add(MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))

```

```

model1.add(Dropout(rate=0.25))

model1.add(Flatten())
model1.add(Dense(units=128, activation='relu'))
model1.add(Dense(units=10, activation='softmax'))
model1.compile(loss=keras.losses.sparse_categorical_crossentropy,
                optimizer=keras.optimizers.Adam(),
                metrics=['accuracy'])
model1.summary()

```

Model: "sequential_1"

| Layer (type) | Output Shape | |
|--------------------------------|--------------------|--|
| ↳ Param # | | |
| conv2d_3 (Conv2D) | (None, 26, 26, 32) | |
| ↳ 320 | | |
| max_pooling2d_3 (MaxPooling2D) | (None, 13, 13, 32) | |
| ↳ 0 | | |
| dropout_3 (Dropout) | (None, 13, 13, 32) | |
| ↳ 0 | | |
| conv2d_4 (Conv2D) | (None, 11, 11, 32) | |
| ↳ 9,248 | | |
| max_pooling2d_4 (MaxPooling2D) | (None, 5, 5, 32) | |
| ↳ 0 | | |
| dropout_4 (Dropout) | (None, 5, 5, 32) | |
| ↳ 0 | | |
| flatten_1 (Flatten) | (None, 800) | |
| ↳ 0 | | |
| dense_2 (Dense) | (None, 128) | |
| ↳ 102,528 | | |
| dense_3 (Dense) | (None, 10) | |
| ↳ 1,290 | | |

Total params: 113,386 (442.91 KB)

Trainable params: 113,386 (442.91 KB)

Non-trainable params: 0 (0.00 B)

```
[8]: x_train = np.expand_dims(x_train, -1)
     x_test = np.expand_dims(x_test, -1)
```

```
[9]: history = model1.fit(x_train, y_train, batch_size=256, epochs=10,
                          validation_split=0.2, verbose=1)
```

```
Epoch 1/10
188/188          41s 203ms/step -
accuracy: 0.5993 - loss: 1.1462 - val_accuracy: 0.8108 - val_loss: 0.5147
Epoch 2/10
188/188          41s 217ms/step -
accuracy: 0.8079 - loss: 0.5148 - val_accuracy: 0.8413 - val_loss: 0.4284
Epoch 3/10
188/188          39s 204ms/step -
accuracy: 0.8383 - loss: 0.4484 - val_accuracy: 0.8625 - val_loss: 0.3942
Epoch 4/10
188/188          39s 192ms/step -
accuracy: 0.8504 - loss: 0.4149 - val_accuracy: 0.8709 - val_loss: 0.3641
Epoch 5/10
188/188          37s 194ms/step -
accuracy: 0.8595 - loss: 0.3807 - val_accuracy: 0.8800 - val_loss: 0.3380
Epoch 6/10
188/188          39s 210ms/step -
accuracy: 0.8704 - loss: 0.3583 - val_accuracy: 0.8799 - val_loss: 0.3299
Epoch 7/10
188/188          35s 188ms/step -
accuracy: 0.8722 - loss: 0.3475 - val_accuracy: 0.8856 - val_loss: 0.3178
Epoch 8/10
188/188          39s 178ms/step -
accuracy: 0.8815 - loss: 0.3256 - val_accuracy: 0.8891 - val_loss: 0.3112
Epoch 9/10
188/188          41s 178ms/step -
accuracy: 0.8847 - loss: 0.3159 - val_accuracy: 0.8915 - val_loss: 0.3011
Epoch 10/10
188/188          42s 183ms/step -
accuracy: 0.8839 - loss: 0.3132 - val_accuracy: 0.8953 - val_loss: 0.2880
```

```
[10]: test_loss, test_acc = model1.evaluate(x_test, y_test, verbose=2)
      print(f'\nTest accuracy for 2 Block CNN: {test_acc:.3f}')
```

```
313/313 - 2s - 6ms/step - accuracy: 0.8929 - loss: 0.3017
```


Test accuracy for 2 block CNN: 0.893

4 Conclusion

As Blocks were decreased from 3 to 2, it gave more accuracy. In block 3, it might have lead to overfitting and FashionMNIST Data may be simpler.

FashionMNIST give more accuracy at same no. of blocks than that of CIFAR10 Dataset.